
DevOps Shack

Docker Best Practices

1. Use Official or Verified Images

- Always prefer official images from Docker Hub or trusted sources.
- Check image signatures and vulnerabilities before using.

2. Use Minimal Base Images

- Use lightweight images like **alpine**, **distroless**, or **scratch** to reduce attack surface.
- Avoid bloated images like **ubuntu** or **debian** unless necessary.

3. Pin Image Versions

- Use specific tags (**node:18.16-alpine**) instead of **latest** to ensure consistency.
- Maintain a manifest of approved images.

4. Reduce Image Layers

- Combine **RUN** commands into a single line to avoid unnecessary layers.

Example:

```
RUN apt-get update && apt-get install -y curl && rm -rf /var/lib/apt/lists/*
```

5. Leverage Multi-Stage Builds



- Reduce final image size by discarding build-time dependencies.

Example:

```
FROM golang:1.19 AS builder
```

```
WORKDIR /app
```

```
COPY . .
```

```
RUN go build -o myapp
```

```
FROM alpine
```

```
COPY --from=builder /app/myapp /usr/local/bin/myapp
```

```
CMD ["myapp"]
```

6. Use `.dockerignore` Efficiently

- Exclude unnecessary files (e.g., `node_modules`, `.git`, logs).

Example `.dockerignore`:

```
.git
```

```
node_modules
```

```
logs
```

7. Set a Non-Root User

- Avoid running containers as `root` to reduce security risks.

Example:



```
RUN adduser --disabled-password myuser  
USER myuser
```

8. Keep Containers Stateless

- Store data externally (e.g., volumes, databases).
- Do not rely on local container storage.

9. Use Environment Variables for Configurations

- Avoid hardcoding secrets; use **ENV** or secrets management.

Example:

```
ENV APP_ENV=production
```

10. Limit Container Resources

- Use **--memory** and **--cpus** flags to prevent resource hogging.

Example:

```
docker run --memory="512m" --cpus="1" myapp
```

11. Optimize Layer Caching

- Place frequently changing lines (e.g., **COPY** commands) at the bottom of the to improve caching efficiency.

12. Keep Images Updated

- Regularly update images and rebuild to include security patches.



Example:

```
docker pull nginx:alpine && docker run nginx:alpine
```

13. Scan Images for Vulnerabilities

- Use tools like:
 - Trivy (**trivy image myimage**)
 - Docker Scout
 - Snyk
 - Anchore

14. Use Read-Only Filesystem When Possible

- Reduce attack vectors by making the filesystem immutable.

Example:

```
docker run --read-only myimage
```

15. Run Only One Process per Container

- Follow the single-responsibility principle (e.g., don't run both Nginx and MySQL in one container).
- Use multiple services via Docker Compose.

16. Use Multi-Arch Images

- Ensure compatibility across different platforms (**amd64, arm64**).

Example:



```
docker buildx build --platform linux/amd64,linux/arm64  
-t myapp:latest .
```

17. Implement Proper Logging

- Redirect logs to **stdout** and **stderr** for containerized logging.

Example:

```
docker logs mycontainer
```

18. Use Health Checks

- Define **HEALTHCHECK** in to ensure container reliability.

Example:

```
HEALTHCHECK --interval=30s --timeout=5s --retries=3 CMD  
curl -f http://localhost || exit 1
```

19. Use Labels for Metadata

- Add metadata for maintainability.

Example:

```
LABEL maintainer="admin@mycompany.com"  
LABEL version="1.0.0"
```

20. Enable Container Restart Policies



- Use **--restart** flag to define behavior on crashes.

Example:

```
docker run --restart unless-stopped myapp
```

21. Use Private Registries for Internal Images

- Use Docker Hub, AWS ECR, GCP Artifact Registry, or Harbor for private images.

Authenticate via:

```
docker login myregistry.com
```

22. Keep Containers Ephemeral

- Containers could be disposable; use volumes for persistent data.
- Avoid modifying running containers.

23. Use Volume Mounts for Persistent Data

- Prefer **-v** volumes over bind mounts for better performance.

Example:

```
docker run -v myvolume:/data myapp
```

24. Ensure Port Binding Security

- Avoid exposing unnecessary ports.



Use explicit port binding:

```
docker run -p 8080:80 myapp
```

25. Limit Container Privileges

- Run with **--cap-drop=ALL** and add only required capabilities.

Example:

```
docker run --cap-drop=ALL --cap-add=NET_ADMIN myapp
```

26. Avoid **ADD** When **COPY** Is Sufficient

- **COPY** is preferred unless you need URL extraction or tarball auto-extraction.

27. Keep Secrets Secure

- Use **--secret** flag or a secrets manager.
- Avoid adding secrets in **or ENV**.

28. Use CI/CD for Automated Builds

- Automate builds and deployments with GitHub Actions, GitLab CI/CD, Jenkins.

29. Secure the Docker Daemon

- Use TLS to secure the Docker API.
- Restrict access using **iptables** or **firewalld**.



30. Regularly Prune Unused Containers and Images

- Clean up old images and containers to free disk space.

Example:

```
docker system prune -af
```

31. Validate Configurations with Linter

- Use **hadolint** to check for best practices.

Example:

```
hadolint
```

32. Use **docker-compose** for Multi-Service Apps

- Maintain **docker-compose.yml** for better management.

Example:

```
version: "3"
services:
  app:
    image: myapp
    ports:
      - "8080:80"
```

33. Consider Rootless Docker for Security



- Running Docker in rootless mode reduces privilege risks.

34. Use Namespaces for Isolation

- Utilize cgroups and user namespaces to restrict access.

35. Set Up Network Security

- Use Docker networks instead of **--host** networking.

Example:

```
docker network create mynetwork  
docker run --network=mynetwork myapp
```

36. Use Immutable Tags for Deployment

- Instead of **latest**, use commit-based tags for reproducible builds.

Example:

```
docker build -t myapp:v1.2.3 .
```

37. Enable AppArmor or SELinux Policies

- Use security profiles to restrict access to resources.

Example (AppArmor):

```
docker run --security-opt apparmor=myprofile myapp
```

38. Avoid Privileged Mode



- Never run containers with **--privileged** unless absolutely necessary.
- It grants full host access, which is a major security risk.

39. Use Host Networking Only When Necessary

- **--network=host** bypasses Docker's network isolation.
- Prefer bridge or overlay networks instead.

40. Restrict Container Process Capabilities

- Drop all unnecessary Linux capabilities using **--cap-drop=ALL**.

Example:

```
docker run --cap-drop=ALL --cap-add=NET_ADMIN myapp
```

41. Use Health Check Retries to Avoid False Positives

- Increase **--retries** value to ensure accurate health status.

Example:

```
HEALTHCHECK --interval=10s --timeout=5s --retries=5 CMD  
curl -f http://localhost || exit 1
```

42. Use tmpfs for Temporary Data

- Instead of writing to disk, use **--tmpfs** for temporary storage.

Example:



```
docker run --tmpfs /tmp:size=100m myapp
```

43. Avoid `--pid=host` for Security

- aring the host's PID namespace increases risks of process injection.
- Use default process isolation.

44. Regularly Audit and Configs

- Automate security checks using tools like:
 - Docker Bench for Security (`docker bench security`)
 - Trivy, Anchore, or Snyk for scanning

45. Ensure Logs Are Managed Properly

- Avoid letting logs grow indefinitely.

Use log rotation:

```
docker run --log-driver=json-file --log-opt  
max-size=10m --log-opt max-file=3 myapp
```

46. Enforce Image Provenance & Signature Verification

- Use Docker Content Trust (DCT) to sign and verify images.

Example:

```
export DOCKER_CONTENT_TRUST=1  
docker pull mysecureimage
```

47. Use CI/CD Pipelines for Automated Security Scans



- Integrate security scanning into Jenkins, GitHub Actions, GitLab CI/CD.

Example GitHub Action:

```
jobs:
  scan:
    runs-on: ubuntu-latest
    steps:
      - uses: aquasecurity/trivy-action@master
        with:
          image-ref: 'myapp:latest'
```

48. Prefer Build-Time Arguments Over Environment Variables

- Use **ARG** instead of **ENV** to avoid persisting sensitive values in the final image.

Example:

```
ARG API_KEY
RUN echo "API_KEY=${API_KEY}"
```

49. Avoid Mounting the Docker Socket (**/var/run/docker.sock**)

- Exposing the Docker daemon socket to containers can lead to privilege escalation.
- If necessary, use tools like gVisor or sysbox for isolation.



50. Enable Seccomp Profiles

- Reduce attack surface by enforcing syscall restrictions.
- Example:

```
docker run --security-opt seccomp=default.json  
myapp
```