



HIGH AVAILABILITY IN DEVOPS PROJECTS

BY DEVOPS SHACK

[Click here for DevSecOps & Cloud DevOps Course](#)

DevOps Shack

High Availability in DevOps Projects –

What Recruiters Love to Hear

● Basic – *"HA Web Server with NGINX & Keepalived on AWS EC2"*

- Build a basic two-node web server cluster with automatic failover using Keepalived for VIP management.

● Intermediate – *"Highly Available Jenkins CI/CD with NFS Shared Storage on Kubernetes"*

- Deploy Jenkins in HA mode with multiple replicas, shared NFS volume for /var/jenkins_home, and Kubernetes Horizontal Pod Autoscaling.

● Advanced – *"Production-Grade HA Microservices Stack with Istio, Vault, MongoDB, and ArgoCD on EKS"*

- Full-scale project deploying multi-tier apps with HA MongoDB StatefulSet, HA Vault with Raft, Canary rollout using ArgoCD, and Istio-managed routing.

■ High Availability in DevOps Projects

What Recruiters Love to Hear — From Basics to Production-Grade Architectures

In today's digital-first world, **downtime is expensive** — every second of unavailability can cost businesses users, revenue, and trust. That's why **High Availability (HA)** is no longer optional. It's a **core skill** every serious DevOps engineer must master.

But here's the twist: while many talk about HA, **few know how to design and implement it practically**. This guide fills that gap with **real-world DevOps projects** that showcase your ability to build resilient, fault-tolerant, self-healing systems.

Whether you're a beginner just getting started, an intermediate learner scaling your skills, or an advanced engineer prepping for top-tier interviews — this guide walks you through **3 real HA projects**, each one designed to:

- ☒ Deepen your technical architecture knowledge
- ☒ Help you stand out to recruiters with real proof of HA thinking
- ☒ Build confidence in designing, deploying, and explaining fault-tolerant systems

🎯 Why It Matters:

High Availability (HA) isn't just a buzzword — it's **business-critical**.

Recruiters **love candidates** who:

- ☒ Think in terms of **fault tolerance**, not just features
- ☒ Architect for **99.99% uptime** or better
- ☒ Proactively **design recovery**, not just deployment

Whether you're applying for a DevOps, SRE, or Platform Engineer role — the **ability to think and build for resilience** is a skill that **sets you apart**.

Let's break it down into what you can **showcase in interviews, resumes, or projects** to highlight HA experience:

🧠 What Is High Availability (HA)?

High Availability means your systems remain **operational and accessible** even when:

-
- A pod or container crashes
 - A node goes down
 - A region becomes unavailable
 - A disk fails or IOPS drops
 - A service gets overwhelmed with traffic

In technical terms, HA requires:

- **Redundancy:** Multiple instances of critical components
- **Load Balancing:** Traffic distribution to healthy nodes
- **Failover Mechanisms:** Automated switching to backup systems
- **Health Checks:** Probes to detect and react to failures
- **State Management:** Persistent data storage that survives failure

Recruiters don't just want to know that you can *run* apps—they want to see that you can keep them running **no matter what**.

☒ Project 1: High Availability Web Server with NGINX & Keepalived on AWS EC2

Level: Basic

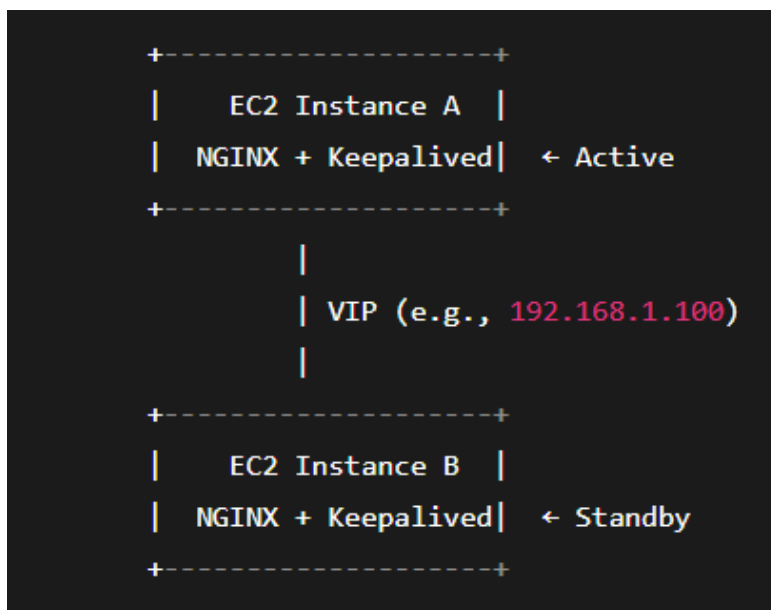
Objective:

Deploy two NGINX web servers behind a **Virtual IP (VIP)** using **Keepalived** to ensure **automatic failover** when one EC2 instance goes down.

What You'll Learn

- Basic High Availability concepts using *active-passive setup*
- How to install and configure **NGINX**
- How to configure **Keepalived** to manage a **floating IP**
- Testing automatic failover when a node goes down

Architecture Overview



Note: Only one instance holds the VIP at a time. If the active one fails, the passive node becomes active.

Prerequisites

- 2 AWS EC2 instances (Ubuntu 22.04 recommended) in the same VPC and subnet

- A **secondary private IP address** in the subnet to act as a floating VIP
- SSH access with sudo privileges to both instances
- Security Group allowing:
 - TCP port 80 (HTTP)
 - Protocol 112 (VRRP)
 - ICMP (for health check pings)

◇ Step 1: Allocate a Secondary Private IP (VIP)

1. Go to **EC2 Dashboard** → **Network Interfaces**
2. Find the network interface attached to **Instance A**
3. Click **Actions** → **Manage IP Addresses** → **Assign new IP**
4. Note down the new private IP (e.g., 192.168.1.100)
5. **Do NOT associate this with any instance manually** – Keepalived will do this dynamically.

◇ Step 2: Install NGINX & Keepalived

Run the following on **both instances**:

```
sudo apt update
```

```
sudo apt install -y nginx keepalived net-tools
```

Check installation:

```
nginx -v
```

```
keepalived -v
```

◇ Step 3: Configure NGINX

On **both nodes**, replace the default NGINX index page:

```
echo "Welcome to Instance A" | sudo tee /var/www/html/index.html
```

Or on B:

```
echo "Welcome to Instance B" | sudo tee /var/www/html/index.html
```

Restart NGINX:

```
sudo systemctl restart nginx
```

Check on browser: http://<EC2_IP>

◆ Step 4: Configure Keepalived on Instance A (MASTER)

Create the config file:

```
sudo nano /etc/keepalived/keepalived.conf
```

Paste:

```
vrrp_instance VI_1 {  
    state MASTER  
  
    interface eth0  
  
    virtual_router_id 51  
  
    priority 101  
  
    advert_int 1  
  
    authentication {  
        auth_type PASS  
        auth_pass devops123  
    }  
  
    virtual_ipaddress {  
        192.168.1.100  
    }  
}
```

Save and restart:

```
sudo systemctl restart keepalived
```

Check:

```
ip a | grep 192.168.1.100
```

You should see the VIP bound to eth0 on **Instance A**.

◇ Step 5: Configure Keepalived on Instance B (BACKUP)

Edit config:

```
sudo nano /etc/keepalived/keepalived.conf
```

Paste:

```
vrrp_instance VI_1 {  
    state BACKUP  
  
    interface eth0  
  
    virtual_router_id 51  
  
    priority 100  
  
    advert_int 1  
  
    authentication {  
        auth_type PASS  
        auth_pass devops123  
    }  
  
    virtual_ipaddress {  
        192.168.1.100  
    }  
}
```

Restart service:

```
sudo systemctl restart keepalived
```

Instance B won't show the VIP yet (expected).

◇ Step 6: Test High Availability (Failover)

☒ Check active instance:

```
curl http://192.168.1.100
```

```
# Should show: Welcome to Instance A
```

🔴 Simulate failure:

```
On Instance A:
```

```
sudo systemctl stop keepalived
```

```
On Instance B:
```

```
ip a | grep 192.168.1.100
```

```
# Should now see VIP assigned to eth0
```

```
Test again:
```

```
curl http://192.168.1.100
```

```
# Output: Welcome to Instance B
```

```
🎉 Success! HA failover works.
```

🔒 Security Consideration

- Ensure VRRP (protocol 112) is **allowed in the security group**
- Use stronger auth_pass for production
- Disable root SSH if not required
- Consider health check scripts and notifications for production

📌 What to Mention in Resume or Interview

- ☑ Built an HA web server cluster using NGINX & Keepalived
- ☑ Configured VIP failover between EC2 instances via VRRP protocol
- ☑ Designed Active-Passive topology to simulate real-world disaster recovery
- ☑ Validated system resilience via simulated node failure tests

🚀 Project 2: Highly Available Jenkins CI/CD with NFS Shared Storage on Kubernetes

🔴 Level: Intermediate

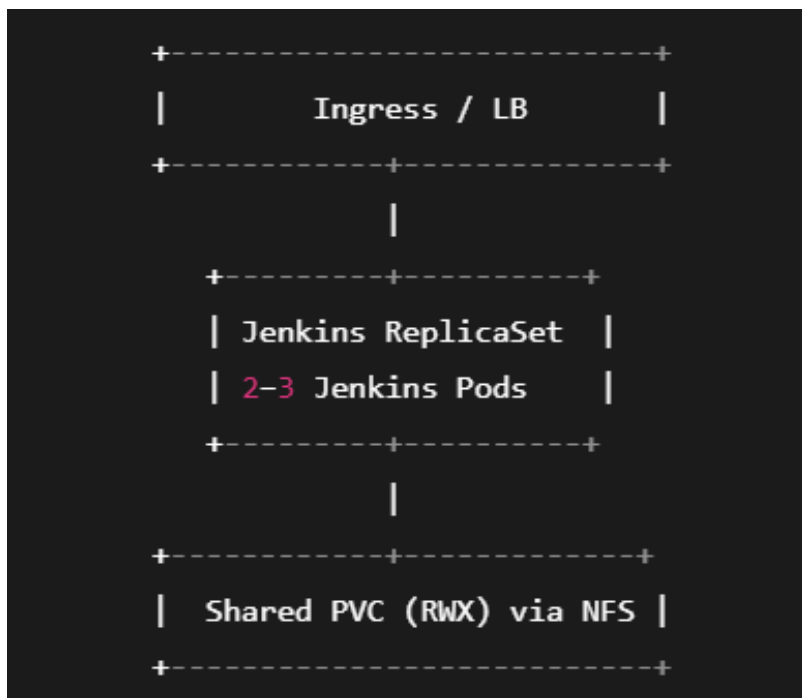
🔗 Objective:

Deploy **Jenkins in HA mode** on Kubernetes using **multiple replicas** backed by a **shared NFS volume**, ensuring that CI/CD pipelines stay available even if a pod crashes or gets rescheduled.

What You'll Learn

- Deploy Jenkins in **High Availability mode** on Kubernetes
- Configure **Persistent Shared Storage** using NFS
- Use **ReadWriteMany (RWX)** volumes to support multiple replicas
- Leverage **Horizontal Pod Autoscaler (HPA)** and **liveness probes**
- Test pod crash recovery and job state persistence

Architecture Overview



Prerequisites

- Kubernetes cluster (minikube or EKS/GKE/AKS)
- NFS Server deployed (or use NFS provisioner)
- Helm installed

- kubectl configured
- Ingress controller (NGINX or ALB Ingress)

◇ Step 1: Setup a Shared NFS Volume (Using NFS Provisioner)

If you don't already have NFS:

```
helm repo add nfs-subdir-external-provisioner https://kubernetes-sigs.github.io/nfs-subdir-external-provisioner/
```

```
helm repo update
```

```
helm install nfs-server nfs-subdir-external-provisioner/nfs-subdir-external-provisioner \
```

```
--set nfs.server=<NFS_SERVER_IP> \
```

```
--set nfs.path=/exported/path \
```

```
--set storageClass.name=nfs-rwx
```

◇ Step 2: Add Bitnami Jenkins Helm Repo

```
helm repo add bitnami https://charts.bitnami.com/bitnami
```

```
helm repo update
```

◇ Step 3: Prepare a jenkins-values.yaml File

controller:

```
replicaCount: 2
```

```
strategyType: RollingUpdate
```

```
adminPassword: "admin123"
```

persistence:

```
enabled: true
```

storageClass: "nfs-rwx"

accessMode: ReadWriteMany

size: 10Gi

healthProbes:

startupProbe:

enabled: true

initialDelaySeconds: 30

periodSeconds: 10

failureThreshold: 12

livenessProbe:

enabled: true

initialDelaySeconds: 90

periodSeconds: 10

failureThreshold: 6

service:

type: ClusterIP

ingress:

enabled: true

ingressClassName: "nginx"

hostname: jenkins.local

tls: false

◇ Step 4: Deploy Jenkins with Helm

```
helm install jenkins bitnami/jenkins -f jenkins-values.yaml -n jenkins --create-namespace
```

◇ Step 5: Verify Jenkins Is Running in HA

```
kubectl get pods -n jenkins
```

You should see **two Jenkins pods** running.

Check logs for successful startup:

```
kubectl logs <jenkins-pod-name> -n jenkins
```

◇ Step 6: Access Jenkins UI

Update /etc/hosts:

```
<INGRESS_IP> jenkins.local
```

Then access: <http://jenkins.local>

Login with:

User: user

Password: admin123

◇ Step 7: Create a Sample Pipeline and Validate HA

1. Create a simple **Freestyle Job** or **Pipeline Job**
2. Trigger it
3. While the job is running, **delete the active pod**:

```
kubectl delete pod <jenkins-0> -n jenkins
```

- ☒ Job should continue in the next pod
- ☒ UI remains accessible
- ☒ Pipeline state persists (thanks to shared NFS volume)

◇ Step 8: Enable Horizontal Pod Autoscaling (Optional)

```
kubectl autoscale deployment jenkins -n jenkins --cpu-percent=60 --min=2 --max=4
```

You can simulate CPU load to see HPA in action.

■ What to Mention in Resume or Interview

- ☒ Deployed Jenkins in High Availability mode using Kubernetes & NFS
- ☒ Ensured zero downtime CI/CD using multiple controller replicas and RWX PVC
- ☒ Configured liveness/startup probes for fault detection and HPA for scaling
- ☒ Simulated pod failure to validate pipeline continuity and HA setup

💧 Project 3: Production-Grade HA Microservices Stack with Istio, Vault, MongoDB & ArgoCD on EKS

● Level: Advanced

🔗 Objective:

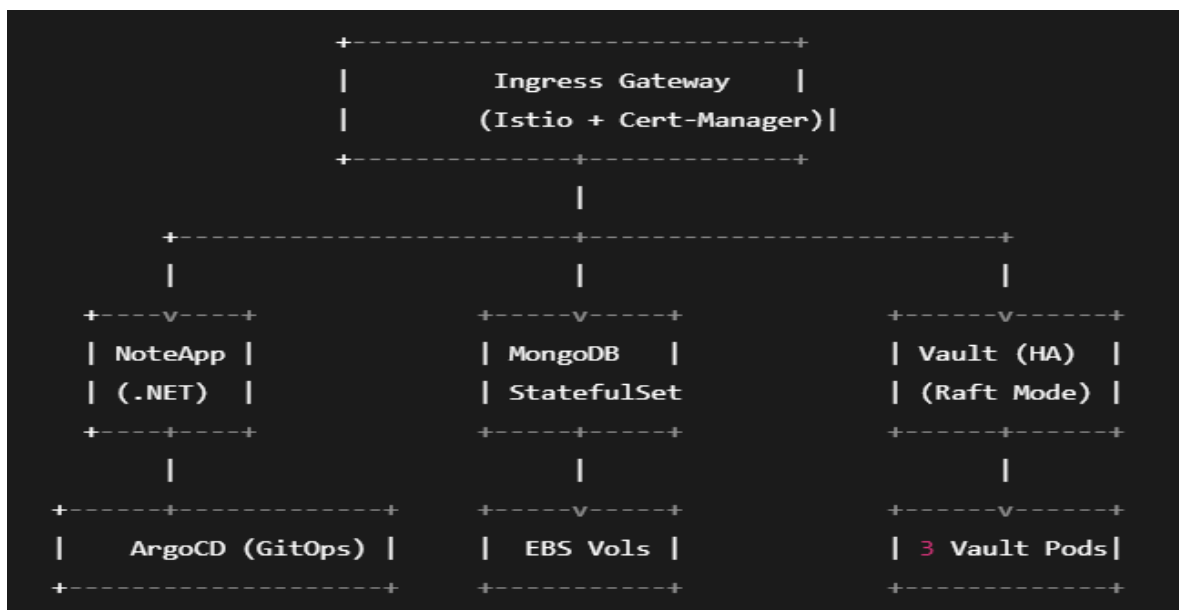
Design and deploy a **highly available microservices platform** on **AWS EKS**, implementing:

- HA MongoDB StatefulSets with data replication
- HA Vault setup with Integrated Raft storage
- Progressive Delivery via ArgoCD (with Canary rollouts)
- Traffic management and observability via Istio

🧠 What You'll Learn

- Build a real-world **microservices architecture** on EKS
- Configure HA for databases (MongoDB), secrets (Vault), and workloads
- Automate GitOps deployments using ArgoCD
- Enable **secure service-to-service communication** via Istio mTLS
- Validate high availability via failover simulations and health checks

🏗️ High-Level Architecture



🔑 Prerequisites

- EKS Cluster (3 or more nodes across multiple AZs)
- Helm & kubectl installed
- Domain name & Route53 (for Ingress DNS)
- SSL certificate via cert-manager
- ArgoCD, Istio, and Vault Helm charts

◇ Step 1: Deploy Istio for Ingress & mTLS

```
istioctl install --set profile=demo -y
```

```
kubectl label namespace default istio-injection=enabled
```

Deploy Istio Gateway and VirtualService to expose your app:

```
# ingress-gateway.yaml
```

```
apiVersion: networking.istio.io/v1beta1
```

```
kind: Gateway
```

```
...
```

Enable **mTLS** and **DestinationRules** to handle traffic securely.

◇ Step 2: Deploy MongoDB as HA StatefulSet

To ensure **High Availability**, we will deploy MongoDB as a **StatefulSet** with:

- **3 replicas**
- **Pod-specific persistent volumes** using EBS
- **Headless service** for internal DNS
- **Readiness and liveness probes**
- **ReplicaSet configuration** for automatic failover

📁 File 1: mongo-headless-service.yaml

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: mongo
```

```
  namespace: app
```

```
spec:
```

ports:

- port: 27017

clusterIP: None # headless

selector:

app: mongo

File 2: mongo-statefulset.yaml

apiVersion: apps/v1

kind: StatefulSet

metadata:

name: mongo

namespace: app

spec:

serviceName: "mongo"

replicas: 3

selector:

matchLabels:

app: mongo

template:

metadata:

labels:

app: mongo

spec:

containers:

- name: mongo

image: mongo:4.4

ports:

- containerPort: 27017

volumeMounts:

- name: mongo-persistent-storage
mountPath: /data/db

livenessProbe:

exec:

command:

- mongo
- --eval
- db.adminCommand('ping')

initialDelaySeconds: 30

periodSeconds: 10

readinessProbe:

exec:

command:

- mongo
- --eval
- db.adminCommand('ping')

initialDelaySeconds: 10

periodSeconds: 10

volumeClaimTemplates:

- metadata:

name: mongo-persistent-storage

spec:

accessModes: ["ReadWriteOnce"]

```
storageClassName: ebs-sc
```

```
resources:
```

```
requests:
```

```
storage: 5Gi
```

File 3: mongo-init-config.yaml (for ReplicaSet Setup)

You need to **initiate the ReplicaSet** after the pods are running:

```
kubectl exec -it mongo-0 -n app -- mongosh
```

Then run:

```
js
```

```
CopyEdit
```

```
rs.initiate({  
  _id: "rs0",  
  members: [  
    { _id: 0, host: "mongo-0.mongo.app.svc.cluster.local:27017" },  
    { _id: 1, host: "mongo-1.mongo.app.svc.cluster.local:27017" },  
    { _id: 2, host: "mongo-2.mongo.app.svc.cluster.local:27017" }  
  ]  
})
```

Verify ReplicaSet Health

```
kubectl exec -it mongo-0 -n app -- mongosh
```

```
rs.status()
```

Look for:

- One **PRIMARY**
- Two **SECONDARY**

- health: 1 for all members

HA Testing

1. Kill the PRIMARY pod:

`kubectl delete pod mongo-0 -n app`

2. Run `rs.status()` again from mongo-1 or mongo-2.

Also create:

- mongo-headless service (ClusterIP: None)
- Readiness & Liveness probes
- Affinity rules to spread pods across AZs

Use:

`helm install mongodb percona/psmdb-db --namespace app --values values.yaml`

◇ Step 3: Deploy Vault in HA Mode (Raft)

`helm repo add hashicorp https://helm.releases.hashicorp.com`

`helm repo update`

`helm install vault hashicorp/vault -n vault --create-namespace -f vault-values.yaml`

Vault config (vault-values.yaml):

`server:`

`ha:`

`enabled: true`

`raft:`

`enabled: true`

`dataStorage:`

enabled: true

storageClass: ebs-sc

size: 10Gi

injector:

enabled: true

Initialize and unseal Vault:

kubectl exec -it vault-0 -n vault -- vault operator init

◇ Step 4: Deploy Microservice (e.g., NoteApp)

Deployment uses annotations for Vault injector

annotations:

vault.hashicorp.com/agent-inject: "true"

vault.hashicorp.com/role: "noteapp"

vault.hashicorp.com/agent-inject-secret-MongoDB__ConnectionString:
"secret/data/noteapp"

ServiceAccount, Vault policy, and Kubernetes auth setup required.

Deploy app with:

kubectl apply -f noteapp-deployment.yaml

◇ Step 5: Setup ArgoCD for GitOps + Canary

kubectl create namespace argocd

helm install argocd argo/argo-cd -n argocd

Create Application YAML pointing to GitHub repo.

Set up **Progressive Delivery**:

canary:

steps:

- setWeight: 20

- pause: {}
- setWeight: 50
- pause: { duration: 30s }

Step 6: Configure Ingress + TLS

✓ Objective:

Secure external access to your NoteApp with:

- Istio Ingress Gateway
- DNS domain mapped to LoadBalancer
- TLS certificates managed by cert-manager

🔗 Prerequisites:

- Istio installed and configured
- cert-manager installed
Install it via Helm if not already:

```
helm repo add jetstack https://charts.jetstack.io
```

```
helm repo update
```

```
helm install cert-manager jetstack/cert-manager \
```

```
--namespace cert-manager --create-namespace \
```

```
--set installCRDs=true
```

- A **domain name** (e.g., noteapp.example.com)
- A valid **email address** (used by Let's Encrypt)

📁 Step A: Create a ClusterIssuer (Let's Encrypt)

```
# cluster-issuer.yaml
```

```
apiVersion: cert-manager.io/v1
```

```
kind: ClusterIssuer
```

metadata:

name: letsencrypt-prod

spec:

acme:

email: your-email@example.com

server: https://acme-v02.api.letsencrypt.org/directory

privateKeySecretRef:

name: letsencrypt-prod

solvers:

- http01:

ingress:

class: istio

Apply it:

kubectl apply -f cluster-issuer.yaml

Step B: Create Certificate Resource

certificate.yaml

apiVersion: cert-manager.io/v1

kind: Certificate

metadata:

name: noteapp-cert

namespace: app

spec:

secretName: noteapp-tls

issuerRef:

name: letsencrypt-prod

```
kind: ClusterIssuer
```

```
commonName: noteapp.example.com
```

```
dnsNames:
```

```
- noteapp.example.com
```

Apply it:

```
kubectl apply -f certificate.yaml
```

This will automatically generate a TLS certificate and store it in the noteapp-tls Kubernetes secret.

Step C: Configure Istio Gateway & VirtualService

```
# istio-gateway.yaml
```

```
apiVersion: networking.istio.io/v1beta1
```

```
kind: Gateway
```

```
metadata:
```

```
  name: noteapp-gateway
```

```
  namespace: app
```

```
spec:
```

```
  selector:
```

```
    istio: ingressgateway
```

```
  servers:
```

```
    - port:
```

```
      number: 443
```

```
      name: https
```

```
      protocol: HTTPS
```

```
    tls:
```

```
      mode: SIMPLE
```

```
      credentialName: noteapp-tls
```

hosts:

- noteapp.example.com

istio-virtualservice.yaml

apiVersion: networking.istio.io/v1beta1

kind: VirtualService

metadata:

- name: noteapp-vs

- namespace: app

spec:

hosts:

- noteapp.example.com

gateways:

- noteapp-gateway

http:

- match:

- uri:

- prefix: /

route:

- destination:

- host: noteapp

- port:

- number: 80

Apply them:

bash

CopyEdit

```
kubectl apply -f istio-gateway.yaml
```

```
kubectl apply -f istio-virtualservice.yaml
```

Step D: DNS Configuration

1. Get **LoadBalancer IP** of Istio Gateway:

```
kubectl get svc istio-ingressgateway -n istio-system
```

2. Add an **A record** in your DNS provider (Route53, GoDaddy, etc.):

noteapp.example.com → <LoadBalancer IP>

Wait for DNS propagation (usually <5 min).

☒ Test Secure Access

Visit:

<https://noteapp.example.com>

- ✓ Should load the NoteApp securely
- ✓ TLS padlock visible
- ✓ Certificate issued by Let's Encrypt

What to Mention in Resume or Interview

- ☒ Designed and implemented a production-grade HA microservices platform on AWS EKS
- ☒ Enabled HA for Vault (Raft), MongoDB (ReplicaSet), and ArgoCD (GitOps)
- ☒ Secured service-to-service communication via Istio mTLS
- ☒ Used StatefulSets, PVs, and Helm charts to ensure self-healing and fault tolerance
- ☒ Applied Progressive Delivery strategy to minimize release risks

☒ Summary of Tools Used

Component	Tool
App Platform	Kubernetes on EKS
HA Ingress & Traffic	Istio
GitOps	ArgoCD
Secrets Management	Vault (HA Raft)
DB with Replication	MongoDB (StatefulSet + PVC)
Storage	AWS EBS
TLS Certs	cert-manager
CI/CD	Jenkins (Optional extension)

Conclusion: Building Resilience

High Availability (HA) isn't just a technical checkbox — it's a **philosophy** that defines how you think, design, and operate systems at scale.

In this guide, you've walked through **three progressively challenging projects** that demonstrate your ability to:

- ☒ Architect resilient systems that **don't break under pressure**
- ☒ Handle failures gracefully using **active-passive setups, replication, and self-healing deployments**
- ☒ Integrate the latest in **GitOps, Service Mesh, StatefulSets, Vault HA**, and more
- ☒ Think like an **SRE, DevOps, or Platform Engineer** who owns not just uptime, but user trust

Whether you're prepping for your first job, switching roles, or trying to stand out in a crowded job market, showing that you've actually **implemented HA in real projects** is a **career-defining advantage**.

What to Do Next

- ✦ Add these projects to your GitHub portfolio
- 📄 Mention HA highlights clearly in your resume
- 💬 Be ready to **explain these architectures** in interviews
- 🎥 Record short video demos or blog posts showcasing failover scenarios
- 🚀 Start treating *availability* as a core design principle in all your future projects